

Minimum Array Partition

Daniil Grbić, Ivan Gogić

September 29, 2023

1 Problem description

Given a $n \times n$ array A of non-negative integers, and a positive integer p , we need to choose $p - 1$ horizontal dividers $0 = h_0 < h_1 < \dots < h_p = n$ and $p - 1$ vertical dividers $0 = v_0 < v_1 < \dots < v_p = n$ partitioning A into p^2 blocks, so as to minimize the following expression:

$$\max_{\substack{1 \leq i \leq p \\ 1 \leq j \leq p}} \sum_{\substack{v_{i-1} \leq x \leq v_i \\ h_{j-1} \leq y \leq h_j}} A[x, y]$$

2 Methods

2.1 Brute Force (BF)

Firstly, we write a brute force algorithm for the purpose of finding optimal solutions for small test cases, which we will use to evaluate our optimisation algorithms.

BF generates all possible partitions of the matrix, and chooses the best one of them. This essentially means that BF will need an unreasonable amount of time to solve all but the smallest of testcases.

The complexity of this approach is $O(n, p) = \binom{n}{p}^2 * p^2$, indeed $\binom{n}{p}^2$ is the total number of partitions, while for each of them we need to check p^2 blocks. We find the sum of each block in constant time, by utilising a partial sum DP matrix.

2.1.1 Cutting optimization

A small optimization of the previous approach involves abandoning some partitions early when it is clear that they will not yield a better solution than the current known best.

Consider this: we know the sum S of the entire matrix M , and we know there are p^2 blocks in total. According to the Pigeonhole principle, at least one of said blocks will have a sum greater or equal to $\frac{S}{p^2}$.

Alternatively, if we have checked $m < p^2$ blocks and their total sum is $S_m < S$, then the maximum value of the remaining blocks will be greater or equal to $\frac{S-S_m}{m-p^2}$.

If at any point this sum happens to become greater than the best solution we have found beforehand, we can safely skip checking the rest of the partition in question. In practice, this speeds up the algorithm 3 to 5 times.

2.2 Genetic Algorithm (GA)

A genetic algorithm (GA) is a p-metaheuristic inspired by the process of natural selection. Genetic algorithms are commonly used to generate high-quality solutions to optimization and search problems by relying on biologically inspired operators such as mutation, crossover and selection.

We use a nearly textbook GA implementation. Nevertheless, given the specificity of the problem, crossover and mutation implementation deserve a closer look.

2.2.1 Crossover

Crossover of two partitions can be done by an arithmetical approach. For simplicity, we will explain the process on a one-dimensional partition. This can be directly applied to two dimensions for the problem in question.

Consider two individuals, A and B , with partitions $[1, 3, 9]$ and $[2, 5, 7]$. Their child C will try to take a middle path in the hopes of bettering their fitness, by taking the average of respective divider pairs: $C = [\frac{1+2}{2}, \frac{3+5}{2}, \frac{9+7}{2}] = [1.5, 4, 8]$. Since divisors have to be whole numbers, we fix C by either rounding the first element down, or rounding it up (chosen randomly).

2.2.2 Mutation

Mutation is achieved by trying to make small movements of individual dividers with a given probability to do so, considering there is an available place to move.

2.3 Simulated Annealing (SA)

Simulated Annealing is a probabilistic optimization method inspired by annealing in metallurgy. It's effective for combinatorial optimization problems

with large search spaces. SA starts with an initial solution and uses a temperature schedule to accept worse solutions with a probability.

During the SA process, small changes are introduced through random movements of block borders, allowing for exploration of the solution space.

2.4 Variable Neighbourhood Search (VNS)

Variable Neighborhood Search (VNS) is an optimization technique that iteratively explores different neighborhoods of solutions. It begins with an initial solution and systematically varies the neighborhood structure to search for better solutions.

During the VNS process, changes are introduced through random movements of (multiple) block borders.

3 Method evaluation

3.1 Test generation

We wrote a helper script for generating testcases. The script takes 3 command line arguments: n , p , and $type$; where $type$ refers to how the matrix M is generated:

Ones fills M with ones, mainly used for debugging

Random fills the matrix with numbers from $(1, n/2)$

Linear , $M_{i,j} = i + j + 1$

Squared , $M_{i,j} = i^2 + j^2 + 1$

3.2 Results

All 4 methods are evaluated on a set of 10 representative testcases. Each algorithm is run 5 times on each problem instance, and we record the best, worst, and average results and execution times.

Let us begin with testcases representing the *random* type. For this type of problems, we can see that SA finds close to optimal solutions at a fraction of the time:

```
Testcase: tests/random_n50_p3.in
BRUTE FORCE
  Results : BEST=3769          WORST=3769          AVG=3769.00
  Time (s) : BEST=2.35         WORST=2.55         AVG=2.45
GENETIC ALGORITHM (GA) [pop_size=50, num_iters=250, elitism_size=10]
  Results : BEST=3769          WORST=3769          AVG=3769.00
```

Time (s) : BEST=0.81	WORST=0.81	AVG=0.81
SIMULATED ANNEALING (SA)		
Results : BEST=3769	WORST=3769	AVG=3769.00
Time (s) : BEST=0.13	WORST=0.14	AVG=0.14
VARIABLE NEIGHBOURHOOD SEARCH (VNS) [num_iters=2500]		
Results : BEST=3769	WORST=3769	AVG=3769.00
Time (s) : BEST=0.47	WORST=0.48	AVG=0.47

Testcase: tests/random_n50_p6.in

BRUTE FORCE - skipping due to speed		
GENETIC ALGORITHM (GA) [pop_size=50, num_iters=250, elitism_size=10]		
Results : BEST=1043	WORST=1046	AVG=1045.40
Time (s) : BEST=1.62	WORST=1.68	AVG=1.65
SIMULATED ANNEALING (SA)		
Results : BEST=1043	WORST=1046	AVG=1043.60
Time (s) : BEST=0.97	WORST=1.04	AVG=1.00
VARIABLE NEIGHBOURHOOD SEARCH (VNS) [num_iters=2500]		
Results : BEST=1043	WORST=1046	AVG=1043.60
Time (s) : BEST=3.18	WORST=3.44	AVG=3.28

Testcase: tests/random_n100_p5.in

BRUTE FORCE - skipping due to speed		
GENETIC ALGORITHM (GA) [pop_size=100, num_iters=500, elitism_size=20]		
Results : BEST=11024	WORST=11066	AVG=11033.40
Time (s) : BEST=5.52	WORST=5.71	AVG=5.63
SIMULATED ANNEALING (SA)		
Results : BEST=11024	WORST=11066	AVG=11037.20
Time (s) : BEST=0.60	WORST=0.66	AVG=0.64
VARIABLE NEIGHBOURHOOD SEARCH (VNS) [num_iters=10000]		
Results : BEST=11024	WORST=11066	AVG=11049.20
Time (s) : BEST=9.29	WORST=9.43	AVG=9.39

Next, let us compare the results achieved on the *linear* test type.

Testcase: tests/linear_n20_p4.in

BRUTE FORCE		
Results : BEST=602	WORST=602	AVG=602.00
Time (s) : BEST=4.71	WORST=5.17	AVG=4.99
GENETIC ALGORITHM (GA) [pop_size=20, num_iters=100, elitism_size=4]		
Results : BEST=602	WORST=602	AVG=602.00
Time (s) : BEST=0.17	WORST=0.18	AVG=0.17
SIMULATED ANNEALING (SA)		
Results : BEST=602	WORST=602	AVG=602.00
Time (s) : BEST=0.32	WORST=0.34	AVG=0.33
VARIABLE NEIGHBOURHOOD SEARCH (VNS) [num_iters=400]		
Results : BEST=602	WORST=608	AVG=603.20

Time (s) : BEST=0.18 WORST=0.20 AVG=0.19

Testcase: tests/linear_n30_p4.in
BRUTE FORCE - skipping due to speed
GENETIC ALGORITHM (GA) [pop_size=30, num_iters=150, elitism_size=6]
 Results : BEST=1920 WORST=1968 AVG=1958.40
 Time (s) : BEST=0.37 WORST=0.39 AVG=0.38
SIMULATED ANNEALING (SA)
 Results : BEST=1920 WORST=1920 AVG=1920.00
 Time (s) : BEST=0.32 WORST=0.35 AVG=0.34
VARIABLE NEIGHBOURHOOD SEARCH (VNS) [num_iters=900]
 Results : BEST=1920 WORST=1968 AVG=1939.20
 Time (s) : BEST=0.43 WORST=0.44 AVG=0.43

We find the results for *squared* test the most interesting.

Testcase: tests/squared_n40_p3.in
BRUTE FORCE
 Results : BEST=214662 WORST=214662 AVG=214662.00
 Time (s) : BEST=1.60 WORST=1.72 AVG=1.65
GENETIC ALGORITHM (GA) [pop_size=40, num_iters=200, elitism_size=8]
 Results : BEST=214662 WORST=215757 AVG=215538.00
 Time (s) : BEST=0.52 WORST=0.53 AVG=0.52
SIMULATED ANNEALING (SA)
 Results : BEST=214662 WORST=214662 AVG=214662.00
 Time (s) : BEST=0.13 WORST=0.13 AVG=0.13
VARIABLE NEIGHBOURHOOD SEARCH (VNS) [num_iters=1600]
 Results : BEST=214662 WORST=215757 AVG=215538.00
 Time (s) : BEST=0.30 WORST=0.31 AVG=0.30

Testcase: tests/squared_n40_p4.in
BRUTE FORCE - skipping due to speed
GENETIC ALGORITHM (GA) [pop_size=40, num_iters=200, elitism_size=8]
 Results : BEST=124440 WORST=131775 AVG=125907.00
 Time (s) : BEST=0.66 WORST=0.69 AVG=0.68
SIMULATED ANNEALING (SA)
 Results : BEST=124440 WORST=124440 AVG=124440.00
 Time (s) : BEST=0.31 WORST=0.33 AVG=0.32
VARIABLE NEIGHBOURHOOD SEARCH (VNS) [num_iters=1600]
 Results : BEST=124440 WORST=131775 AVG=125907.00
 Time (s) : BEST=0.76 WORST=0.81 AVG=0.78

Testcase: tests/squared_n100_p10.in
BRUTE FORCE - skipping due to speed
GENETIC ALGORITHM (GA) [pop_size=100, num_iters=500, elitism_size=20]
 Results : BEST=858130 WORST=984543 AVG=941758.40

Time (s) : BEST=14.10	WORST=14.69	AVG=14.47
SIMULATED ANNEALING (SA)		
Results : BEST=903805	WORST=1149356	AVG=1030165.60
Time (s) : BEST=4.56	WORST=4.85	AVG=4.69
VARIABLE NEIGHBOURHOOD SEARCH (VNS) [num_iters=10000]		
Results : BEST=837551	WORST=928305	AVG=868049.20
Time (s) : BEST=28.35	WORST=33.75	AVG=30.67

4 Failed ideas

4.1 Linear Programming

We initially had an idea to solve the problem using linear programming, as it seemed easy to define all the conditions that needed to hold true. The problem arose when defining the cost function, as it required block borders as matrix indices, and it appears that it's not possible to index the array with LP variables.

5 Conclusion

In conclusion, the Minimum Array Partition problem has been established as NP-hard, making it challenging to find optimal solutions efficiently. While there exist polynomial algorithms that can provide approximations within a constant factor, their implementation complexity often outweighs the benefits (our approximations factor is better).

As demonstrated in our study, ...

These methods strike a balance between computational efficiency and solution quality, making them valuable tools for addressing the Minimum Array Partition problem.

6 References

Efficient array partitioning
On the complexity of the generalized block distribution